

Gold Recovery Prediction Project

0. Introduction

This project aims to develop a machine learning model to predict the amount of gold recovered from the ore during the extraction and purification processes. The goal is to optimize the production process for Zyfra, a company focused on improving efficiency in heavy industries, particularly in the gold mining sector.

Project Steps

0. Introduction:

- Outline the goals and steps of the project.

1. Prepare the Data:

- Load and inspect the datasets.
- Verify the accuracy of the recovery calculations.
- Analyze the features that are only available in the training set.
- Perform necessary data preprocessing.

2. Analyze the Data:

- Explore how the concentrations of key metals (Au, Ag, Pb) change through different purification stages.
- Compare the distributions of feed particle sizes in the training and test sets.

3. Develop and Train the Model:

- Split the data into training and testing sets.
- Train machine learning models to predict gold recovery.
- Evaluate the models using appropriate metrics.

4. Model Interpretation:

- Analyze feature importance and provide insights.
- Summarize findings and offer recommendations.

1. Prepare the Data

Open the files and look into the data

In this step, we will load the datasets provided for the project. The datasets contain information about the gold extraction and purification processes. We will inspect the contents of these datasets to understand their structure, identify any potential issues such as missing values, and prepare the data for further analysis.

Path to files:

- `/datasets/gold_recovery_train.csv`
- `/datasets/gold_recovery_test.csv`
- `/datasets/gold_recovery_full.csv`

```
In [1]: # Section 1: Library Imports
import pandas as pd
```

```
In [2]: # Section 2: Data Loading
# Load the datasets (handle errors such as file not found or incorrect paths)
try:
    train_data = pd.read_csv('/datasets/gold_recovery_train.csv')
    test_data = pd.read_csv('/datasets/gold_recovery_test.csv')
    full_data = pd.read_csv('/datasets/gold_recovery_full.csv')
    print("Datasets loaded successfully.")
except FileNotFoundError as e:
    print(f"Error: {e}")
    # Optionally handle or exit the script if files are not found
    raise
```

Datasets loaded successfully.

```
In [3]: # Section 3: Data Inspection
# Inspect the first few rows of each dataset to ensure successful loading
print("Training Data Head:")
print(train_data.head())

print("\nTest Data Head:")
print(test_data.head())

print("\nFull Data Head:")
print(full_data.head())
```

Training Data Head:

	date	final.output.concentrate_ag	\
0	2016-01-15 00:00:00	6.055403	
1	2016-01-15 01:00:00	6.029369	
2	2016-01-15 02:00:00	6.055926	
3	2016-01-15 03:00:00	6.047977	
4	2016-01-15 04:00:00	6.148599	

	final.output.concentrate_pb	final.output.concentrate_sol	\
0	9.889648	5.507324	
1	9.968944	5.257781	
2	10.213995	5.383759	
3	9.977019	4.858634	
4	10.142511	4.939416	

	final.output.concentrate_au	final.output.recovery	final.output.tail_ag	\
0	42.192020	70.541216	10.411962	
1	42.701629	69.266198	10.462676	
2	42.657501	68.116445	10.507046	
3	42.689819	68.347543	10.422762	
4	42.774141	66.927016	10.360302	

	final.output.tail_pb	final.output.tail_sol	final.output.tail_au	...	\
0	0.895447	16.904297	2.143149	...	
1	0.927452	16.634514	2.224930	...	
2	0.953716	16.208849	2.257889	...	
3	0.883763	16.532835	2.146849	...	
4	0.792826	16.525686	2.055292	...	

	secondary_cleaner.state.floatbank4_a_air	\
0	14.016835	
1	13.992281	
2	14.015015	
3	14.036510	
4	14.027298	

	secondary_cleaner.state.floatbank4_a_level	\
0	-502.488007	
1	-505.503262	
2	-502.520901	
3	-500.857308	
4	-499.838632	

	secondary_cleaner.state.floatbank4_b_air	\
0	12.099931	
1	11.950531	
2	11.912783	
3	11.999550	
4	11.953070	

	secondary_cleaner.state.floatbank4_b_level	\
0	-504.715942	
1	-501.331529	
2	-501.133383	
3	-501.193686	
4	-501.053894	

```

secondary_cleaner.state.floatbank5_a_air \
0      9.925633
1     10.039245
2     10.070913
3      9.970366
4      9.925709

```

```

secondary_cleaner.state.floatbank5_a_level \
0    -498.310211
1    -500.169983
2    -500.129135
3    -499.201640
4    -501.686727

```

```

secondary_cleaner.state.floatbank5_b_air \
0      8.079666
1      7.984757
2      8.013877
3      7.977324
4      7.894242

```

```

secondary_cleaner.state.floatbank5_b_level \
0    -500.470978
1    -500.582168
2    -500.517572
3    -500.255908
4    -500.356035

```

```

secondary_cleaner.state.floatbank6_a_air \
0     14.151341
1     13.998353
2     14.028663
3     14.005551
4     13.996647

```

```

secondary_cleaner.state.floatbank6_a_level
0    -605.841980
1    -599.787184
2    -601.427363
3    -599.996129
4    -601.496691

```

[5 rows x 87 columns]

Test Data Head:

```

      date primary_cleaner.input.sulfate \
0  2016-09-01 00:59:59      210.800909
1  2016-09-01 01:59:59      215.392455
2  2016-09-01 02:59:59      215.259946
3  2016-09-01 03:59:59      215.336236
4  2016-09-01 04:59:59      199.099327

```

```

primary_cleaner.input.depressant primary_cleaner.input.feed_size \
0      14.993118      8.080000
1      14.987471      8.080000

```

2	12.884934	7.786667
3	12.006805	7.640000
4	10.682530	7.530000

	primary_cleaner.input.xanthate	primary_cleaner.state.floatbank8_a_air \
0	1.005021	1398.981301
1	0.990469	1398.777912
2	0.996043	1398.493666
3	0.863514	1399.618111
4	0.805575	1401.268123

	primary_cleaner.state.floatbank8_a_level \
0	-500.225577
1	-500.057435
2	-500.868360
3	-498.863574
4	-500.808305

	primary_cleaner.state.floatbank8_b_air \
0	1399.144926
1	1398.055362
2	1398.860436
3	1397.440120
4	1398.128818

	primary_cleaner.state.floatbank8_b_level \
0	-499.919735
1	-499.778182
2	-499.764529
3	-499.211024
4	-499.504543

	primary_cleaner.state.floatbank8_c_air ... \
0	1400.102998 ...
1	1396.151033 ...
2	1398.075709 ...
3	1400.129303 ...
4	1402.172226 ...

	secondary_cleaner.state.floatbank4_a_air \
0	12.023554
1	12.058140
2	11.962366
3	12.033091
4	12.025367

	secondary_cleaner.state.floatbank4_a_level \
0	-497.795834
1	-498.695773
2	-498.767484
3	-498.350935
4	-500.786497

	secondary_cleaner.state.floatbank4_b_air \
0	8.016656
1	8.130979

2	8.096893
3	8.074946
4	8.054678
secondary_cleaner.state.floatbank4_b_level \	
0	-501.289139
1	-499.634209
2	-500.827423
3	-499.474407
4	-500.397500
secondary_cleaner.state.floatbank5_a_air \	
0	7.946562
1	7.958270
2	8.071056
3	7.897085
4	8.107890
secondary_cleaner.state.floatbank5_a_level \	
0	-432.317850
1	-525.839648
2	-500.801673
3	-500.868509
4	-509.526725
secondary_cleaner.state.floatbank5_b_air \	
0	4.872511
1	4.878850
2	4.905125
3	4.931400
4	4.957674
secondary_cleaner.state.floatbank5_b_level \	
0	-500.037437
1	-500.162375
2	-499.828510
3	-499.963623
4	-500.360026
secondary_cleaner.state.floatbank6_a_air \	
0	26.705889
1	25.019940
2	24.994862
3	24.948919
4	25.003331
secondary_cleaner.state.floatbank6_a_level	
0	-499.709414
1	-499.819438
2	-500.622559
3	-498.709987
4	-500.856333

[5 rows x 53 columns]

Full Data Head:

	date	final.output.concentrate_ag	\
0	2016-01-15 00:00:00	6.055403	
1	2016-01-15 01:00:00	6.029369	
2	2016-01-15 02:00:00	6.055926	
3	2016-01-15 03:00:00	6.047977	
4	2016-01-15 04:00:00	6.148599	

	final.output.concentrate_pb	final.output.concentrate_sol	\
0	9.889648	5.507324	
1	9.968944	5.257781	
2	10.213995	5.383759	
3	9.977019	4.858634	
4	10.142511	4.939416	

	final.output.concentrate_au	final.output.recovery	final.output.tail_ag	\
0	42.192020	70.541216	10.411962	
1	42.701629	69.266198	10.462676	
2	42.657501	68.116445	10.507046	
3	42.689819	68.347543	10.422762	
4	42.774141	66.927016	10.360302	

	final.output.tail_pb	final.output.tail_sol	final.output.tail_au	...	\
0	0.895447	16.904297	2.143149	...	
1	0.927452	16.634514	2.224930	...	
2	0.953716	16.208849	2.257889	...	
3	0.883763	16.532835	2.146849	...	
4	0.792826	16.525686	2.055292	...	

	secondary_cleaner.state.floatbank4_a_air	\
0	14.016835	
1	13.992281	
2	14.015015	
3	14.036510	
4	14.027298	

	secondary_cleaner.state.floatbank4_a_level	\
0	-502.488007	
1	-505.503262	
2	-502.520901	
3	-500.857308	
4	-499.838632	

	secondary_cleaner.state.floatbank4_b_air	\
0	12.099931	
1	11.950531	
2	11.912783	
3	11.999550	
4	11.953070	

	secondary_cleaner.state.floatbank4_b_level	\
0	-504.715942	
1	-501.331529	
2	-501.133383	
3	-501.193686	
4	-501.053894	

```

secondary_cleaner.state.floatbank5_a_air \
0          9.925633
1         10.039245
2         10.070913
3          9.970366
4          9.925709

secondary_cleaner.state.floatbank5_a_level \
0        -498.310211
1        -500.169983
2        -500.129135
3        -499.201640
4        -501.686727

secondary_cleaner.state.floatbank5_b_air \
0          8.079666
1          7.984757
2          8.013877
3          7.977324
4          7.894242

secondary_cleaner.state.floatbank5_b_level \
0        -500.470978
1        -500.582168
2        -500.517572
3        -500.255908
4        -500.356035

secondary_cleaner.state.floatbank6_a_air \
0         14.151341
1         13.998353
2         14.028663
3         14.005551
4         13.996647

secondary_cleaner.state.floatbank6_a_level
0        -605.841980
1        -599.787184
2        -601.427363
3        -599.996129
4        -601.496691

```

[5 rows x 87 columns]

```

In [4]: # Section 4: Duplicate Checks
# Check for duplicates in each dataset
train_duplicates = train_data.duplicated().sum()
test_duplicates = test_data.duplicated().sum()
full_duplicates = full_data.duplicated().sum()

# Output the number of duplicates found in each dataset
print(f"\nNumber of duplicate rows in training data: {train_duplicates}")
print(f"Number of duplicate rows in test data: {test_duplicates}")
print(f"Number of duplicate rows in full data: {full_duplicates}")

```


Number of duplicate rows in training data: 0

Number of duplicate rows in test data: 0

Number of duplicate rows in full data: 0

1.2. Check that recovery is calculated correctly

In this step, we will verify the accuracy of the `rougher.output.recovery` feature in the training dataset. The recovery rate is calculated using the following formula:

$$[\text{Recovery}] = \frac{C \times (F - T)}{F \times (C - T)} \times 100\%$$

Where:

- (C) is the concentration of gold in the rougher concentrate after flotation.
- (F) is the concentration of gold in the rougher feed before flotation.
- (T) is the concentration of gold in the rougher tails after flotation.

We will calculate the recovery values using this formula and compare them to the existing values in the `rougher.output.recovery` column. The comparison will be done using the Mean Absolute Error (MAE) metric to measure the difference between our calculations and the provided data.

```
In [5]: from sklearn.metrics import mean_absolute_error

# Check for NaN values in the relevant columns
nan_check = train_data[['rougher.output.concentrate_au',
                        'rougher.input.feed_au',
                        'rougher.output.tail_au',
                        'rougher.output.recovery']].isna().sum()

print("NaN values in the relevant columns:\n", nan_check)

# Drop rows with NaN values in the relevant columns
train_data_cleaned = train_data.dropna(subset=['rougher.output.concentrate_au',
                                                'rougher.input.feed_au',
                                                'rougher.output.tail_au',
                                                'rougher.output.recovery'])

# Calculate the rougher recovery using the provided formula
def calculate_recovery(C, F, T):
    return (C * (F - T) / (F * (C - T))) * 100

# Extract the necessary columns from the cleaned train dataset
C = train_data_cleaned['rougher.output.concentrate_au']
F = train_data_cleaned['rougher.input.feed_au']
T = train_data_cleaned['rougher.output.tail_au']

# Calculate recovery
calculated_recovery = calculate_recovery(C, F, T)
```

```
# Calculate the MAE between the calculated recovery and the actual recovery
mae_recovery = mean_absolute_error(train_data_cleaned['rougher.output.recovery'], c

# Display the calculated MAE
mae_recovery
```

NaN values in the relevant columns:

```
rougher.output.concentrate_au      82
rougher.input.feed_au              83
rougher.output.tail_au             2249
rougher.output.recovery            2573
dtype: int64
```

Out[5]: 9.303415616264301e-15

1.3. Analyze the features not available in the test set

In this step, we will identify the features that are present in the training dataset but not in the test dataset. Understanding these differences is important because the test set should only include features that will be available during prediction, which means any features that are outcomes or results of the process (e.g., calculated recovery values) should be excluded from the test set.

We will examine the names and types of these features to better understand their role and why they might be excluded from the test set.

```
In [6]: # Find the features that are in the training set but not in the test set
train_only_features = set(train_data.columns) - set(test_data.columns)

# Analyze these features by displaying their names and types
train_only_features_info = train_data[list(train_only_features)].dtypes

# Display the features and their types
train_only_features_info
```

```

Out[6]: final.output.concentrate_pb          float64
primary_cleaner.output.concentrate_au       float64
secondary_cleaner.output.tail_pb            float64
primary_cleaner.output.tail_sol              float64
rougher.output.concentrate_au                float64
final.output.concentrate_au                  float64
primary_cleaner.output.tail_ag               float64
primary_cleaner.output.concentrate_pb        float64
rougher.output.tail_au                       float64
rougher.calculation.au_pb_ratio              float64
final.output.concentrate_sol                 float64
rougher.output.tail_pb                       float64
final.output.tail_au                         float64
secondary_cleaner.output.tail_ag             float64
secondary_cleaner.output.tail_au             float64
rougher.output.tail_ag                       float64
primary_cleaner.output.tail_pb               float64
rougher.calculation.sulfate_to_au_concentrate float64
final.output.tail_pb                         float64
primary_cleaner.output.concentrate_sol        float64
final.output.tail_ag                         float64
primary_cleaner.output.concentrate_ag         float64
rougher.output.concentrate_pb                float64
final.output.tail_sol                        float64
secondary_cleaner.output.tail_sol             float64
primary_cleaner.output.tail_au               float64
final.output.concentrate_ag                  float64
rougher.calculation.floatbank10_sulfate_to_au_feed float64
rougher.output.concentrate_ag                float64
final.output.recovery                        float64
rougher.calculation.floatbank11_sulfate_to_au_feed float64
rougher.output.recovery                      float64
rougher.output.tail_sol                      float64
rougher.output.concentrate_sol                float64
dtype: object

```

1.4. Perform Data Preprocessing

In this step, we will perform several key data preprocessing tasks to ensure that the datasets are ready for analysis and modeling. These tasks include:

1. **Handling Missing Values:** We'll decide how to handle missing values in the datasets, whether by imputation or removal of rows/columns with missing data.
2. **Feature Scaling:** We'll apply feature scaling to ensure that all features contribute equally to the model's learning process.
3. **Splitting Data:** We'll prepare the features and target variables in the training set, and ensure that the test set is preprocessed in the same way.

```

In [7]: # Check for missing values in the training and test datasets
missing_values_train = train_data.isnull().sum()
missing_values_test = test_data.isnull().sum()

```

```
# Identify columns with missing values
missing_columns_train = missing_values_train[missing_values_train > 0]
missing_columns_test = missing_values_test[missing_values_test > 0]

print("Missing values in training set:\n", missing_columns_train)
print("\nMissing values in test set:\n", missing_columns_test)

# Drop rows with missing values in the training and test datasets
train_data_cleaned = train_data.dropna()
test_data_cleaned = test_data.dropna()
```

Missing values in training set:

final.output.concentrate_ag	72
final.output.concentrate_pb	72
final.output.concentrate_sol	370
final.output.concentrate_au	71
final.output.recovery	1521
...	
secondary_cleaner.state.floatbank5_a_level	85
secondary_cleaner.state.floatbank5_b_air	85
secondary_cleaner.state.floatbank5_b_level	84
secondary_cleaner.state.floatbank6_a_air	103
secondary_cleaner.state.floatbank6_a_level	85

Length: 85, dtype: int64

Missing values in test set:

primary_cleaner.input.sulfate	302
primary_cleaner.input.depressant	284
primary_cleaner.input.xanthate	166
primary_cleaner.state.floatbank8_a_air	16
primary_cleaner.state.floatbank8_a_level	16
primary_cleaner.state.floatbank8_b_air	16
primary_cleaner.state.floatbank8_b_level	16
primary_cleaner.state.floatbank8_c_air	16
primary_cleaner.state.floatbank8_c_level	16
primary_cleaner.state.floatbank8_d_air	16
primary_cleaner.state.floatbank8_d_level	16
rougher.input.feed_ag	16
rougher.input.feed_pb	16
rougher.input.feed_rate	40
rougher.input.feed_size	22
rougher.input.feed_sol	67
rougher.input.feed_au	16
rougher.input.floatbank10_sulfate	257
rougher.input.floatbank10_xanthate	123
rougher.input.floatbank11_sulfate	55
rougher.input.floatbank11_xanthate	353
rougher.state.floatbank10_a_air	17
rougher.state.floatbank10_a_level	16
rougher.state.floatbank10_b_air	17
rougher.state.floatbank10_b_level	16
rougher.state.floatbank10_c_air	17
rougher.state.floatbank10_c_level	16
rougher.state.floatbank10_d_air	17
rougher.state.floatbank10_d_level	16
rougher.state.floatbank10_e_air	17
rougher.state.floatbank10_e_level	16
rougher.state.floatbank10_f_air	17
rougher.state.floatbank10_f_level	16
secondary_cleaner.state.floatbank2_a_air	20
secondary_cleaner.state.floatbank2_a_level	16
secondary_cleaner.state.floatbank2_b_air	23
secondary_cleaner.state.floatbank2_b_level	16
secondary_cleaner.state.floatbank3_a_air	34
secondary_cleaner.state.floatbank3_a_level	16
secondary_cleaner.state.floatbank3_b_air	16
secondary_cleaner.state.floatbank3_b_level	16

```

secondary_cleaner.state.floatbank4_a_air      16
secondary_cleaner.state.floatbank4_a_level    16
secondary_cleaner.state.floatbank4_b_air      16
secondary_cleaner.state.floatbank4_b_level    16
secondary_cleaner.state.floatbank5_a_air      16
secondary_cleaner.state.floatbank5_a_level    16
secondary_cleaner.state.floatbank5_b_air      16
secondary_cleaner.state.floatbank5_b_level    16
secondary_cleaner.state.floatbank6_a_air      16
secondary_cleaner.state.floatbank6_a_level    16
dtype: int64

```

```

In [8]: from sklearn.preprocessing import StandardScaler

# Identify common features to scale in both train and test datasets
common_features_to_scale = list(set(train_data_cleaned.columns).intersection(set(test_data_cleaned.columns)))

# Ensure working with a copy to avoid SettingWithCopyWarning
train_data_cleaned = train_data_cleaned.copy()
test_data_cleaned = test_data_cleaned.copy()

# Initialize the scaler
scaler = StandardScaler()

# Fit the scaler on the training data and transform the features in both datasets
train_data_cleaned.loc[:, common_features_to_scale] = scaler.fit_transform(train_data_cleaned.loc[:, common_features_to_scale])
test_data_cleaned.loc[:, common_features_to_scale] = scaler.transform(test_data_cleaned.loc[:, common_features_to_scale])

```

```

In [9]: # Define the target variables
target_train = train_data_cleaned[['rougher.output.recovery', 'final.output.recovery']]

# Define the features for the training data
features_train = train_data_cleaned.drop(columns=['rougher.output.recovery', 'final.output.recovery'])

# The test data should have the same features but without the target variables
features_test = test_data_cleaned

```

2.1. Take note of how the concentrations of metals (Au, Ag, Pb) change depending on the purification stage

In this step, we will explore how the concentrations of key metals—gold (Au), silver (Ag), and lead (Pb)—change during the purification stages. This will give us insights into how effective each stage of the process is at concentrating these valuable metals.

The stages of interest are:

- **Rougher stage**
- **Primary cleaner stage**
- **Secondary cleaner stage**

- **Final output stage**

We will plot the concentrations of these metals at each stage to visually assess the changes.

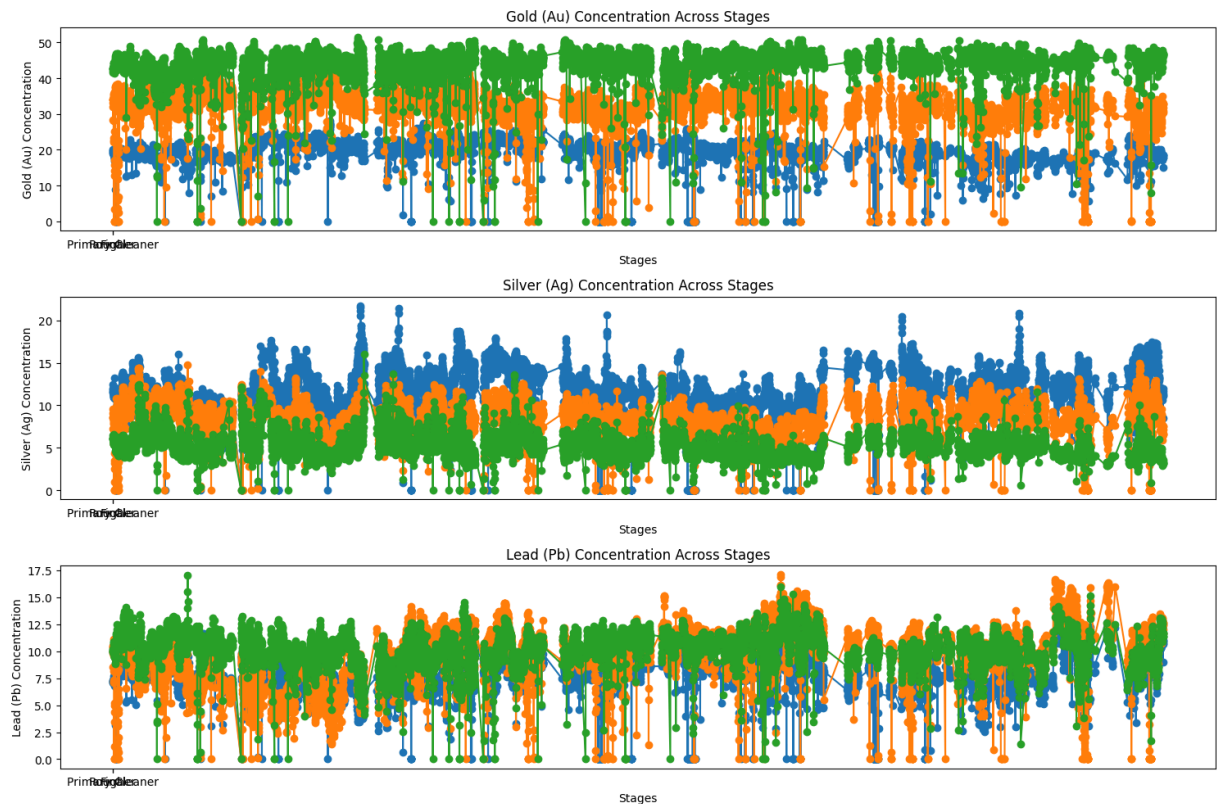
```
In [10]: import matplotlib.pyplot as plt

# Select the columns related to metal concentrations at different stages
concentration_columns = {
    'Gold (Au)': ['rougher.output.concentrate_au',
                  'primary_cleaner.output.concentrate_au',
                  'final.output.concentrate_au'],
    'Silver (Ag)': ['rougher.output.concentrate_ag',
                    'primary_cleaner.output.concentrate_ag',
                    'final.output.concentrate_ag'],
    'Lead (Pb)': ['rougher.output.concentrate_pb',
                  'primary_cleaner.output.concentrate_pb',
                  'final.output.concentrate_pb']
}

# Plot the changes in concentration for each metal
plt.figure(figsize=(15, 10))

for i, (metal, stages) in enumerate(concentration_columns.items(), 1):
    plt.subplot(3, 1, i)
    plt.plot(train_data_cleaned[stages], marker='o')
    plt.title(f'{metal} Concentration Across Stages')
    plt.xlabel('Stages')
    plt.ylabel(f'{metal} Concentration')
    plt.xticks(ticks=[0, 1, 2], labels=['Rougher', 'Primary Cleaner', 'Final'])

plt.tight_layout()
plt.show()
```



2.2. Compare the feed particle size distributions in the training set and in the test set

In this step, we will compare the distributions of feed particle sizes in both the training and test sets. The feed particle size is a critical parameter in the flotation process, and significant differences in these distributions could lead to incorrect model evaluation and poor generalization.

We will visualize the distribution of feed particle sizes in both datasets to assess their similarity.

```
In [11]: import matplotlib.pyplot as plt

# Plotting the feed particle size distributions in the training and test sets using
plt.figure(figsize=(10, 6))

# Plot normalized histogram for training set
plt.hist(train_data_cleaned['rougher.input.feed_size'], bins=30, density=True, alpha

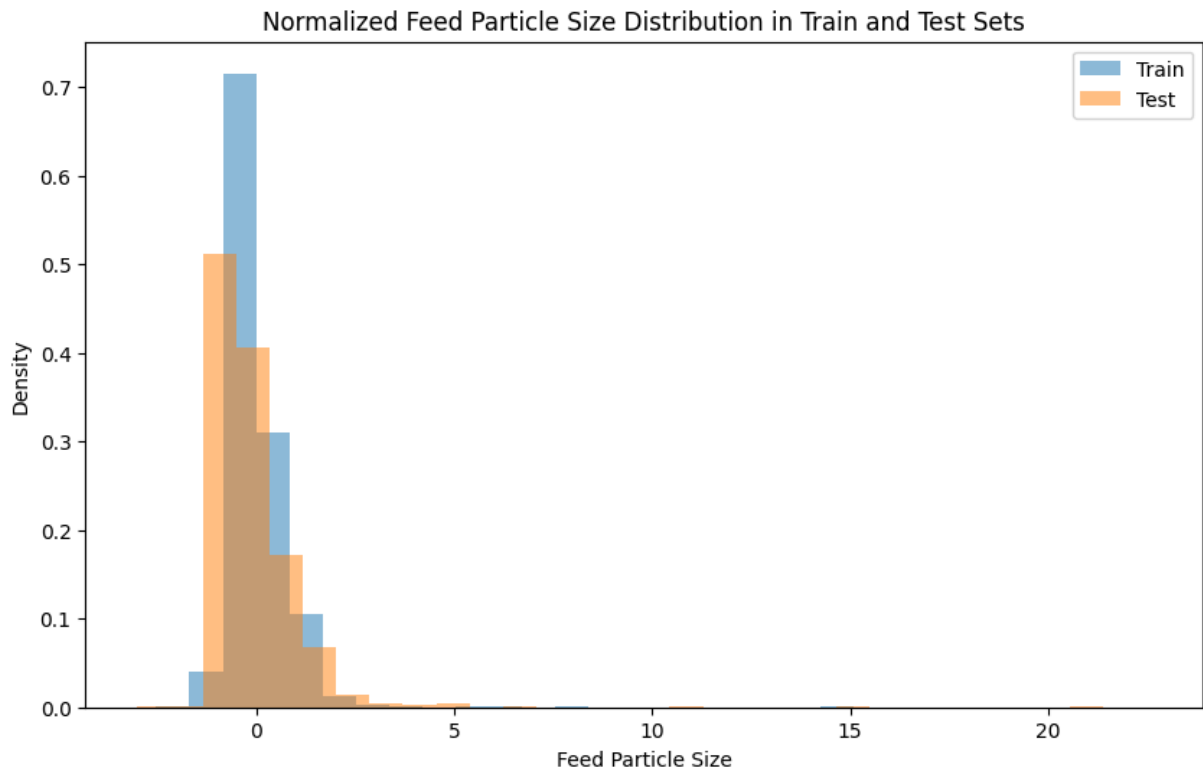
# Plot normalized histogram for test set
plt.hist(test_data_cleaned['rougher.input.feed_size'], bins=30, density=True, alpha

# Add titles and labels
plt.title('Normalized Feed Particle Size Distribution in Train and Test Sets')
plt.xlabel('Feed Particle Size')
```



```
plt.ylabel('Density')
plt.legend()

# Show the plot
plt.show()
```



2.3. Consider the total concentrations of all substances at different stages

In this step, we will calculate the total concentrations of all substances at three key stages:

- **Raw feed:** Before any processing.
- **Rougher concentrate:** After the rougher stage.
- **Final concentrate:** After all purification stages.

We will examine the distributions of these total concentrations to identify any abnormal values. If we find significant outliers, we will consider whether it is worth removing them from both the training and test datasets. These outliers could represent erroneous data or rare events that may negatively impact the model's performance.

After identifying any anomalies, we will eliminate them and clean the data accordingly.

```
In [12]: import matplotlib.pyplot as plt

# Calculate the total concentration of substances at different stages
train_data_cleaned['total_concentration_raw_feed'] = (
```

```
train_data_cleaned['rougher.input.feed_au'] +
train_data_cleaned['rougher.input.feed_ag'] +
train_data_cleaned['rougher.input.feed_pb']
)

train_data_cleaned['total_concentration_rougher'] = (
    train_data_cleaned['rougher.output.concentrate_au'] +
    train_data_cleaned['rougher.output.concentrate_ag'] +
    train_data_cleaned['rougher.output.concentrate_pb']
)

train_data_cleaned['total_concentration_final'] = (
    train_data_cleaned['final.output.concentrate_au'] +
    train_data_cleaned['final.output.concentrate_ag'] +
    train_data_cleaned['final.output.concentrate_pb']
)

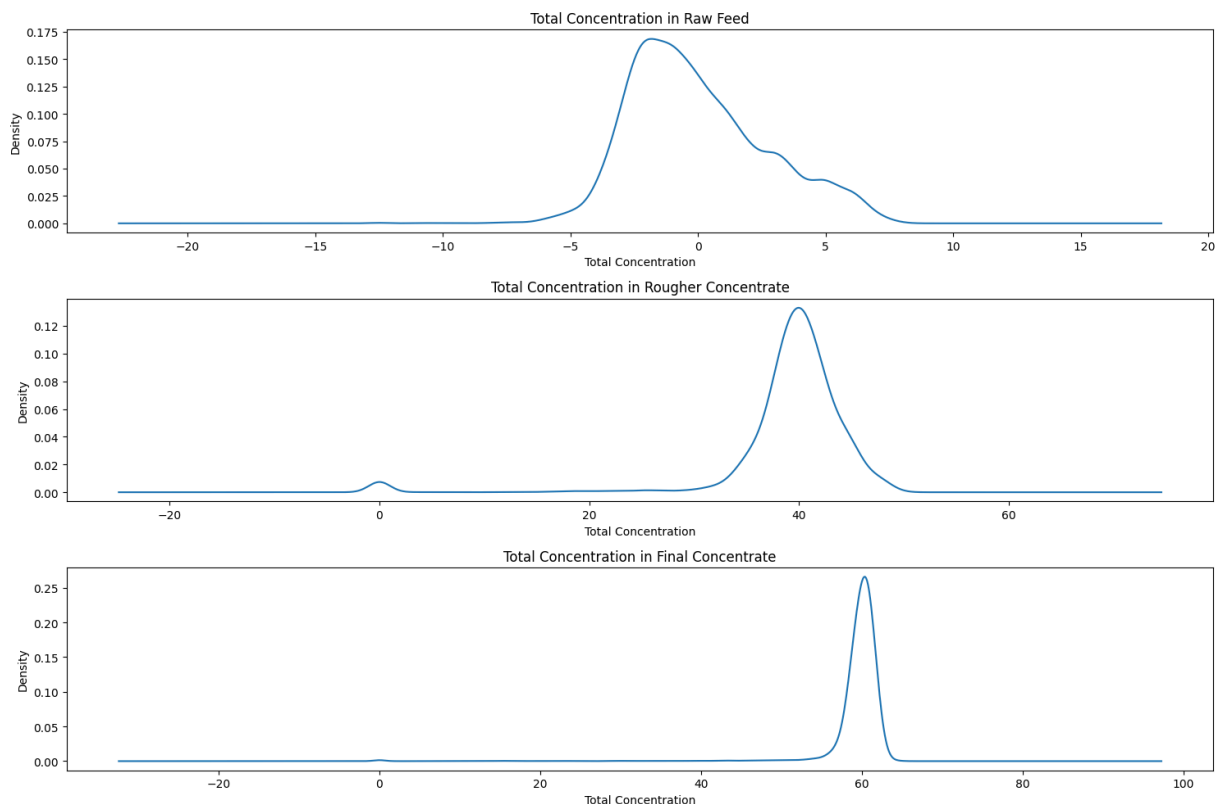
# Plot the distributions of total concentrations at each stage
plt.figure(figsize=(15, 10))

# Raw Feed
plt.subplot(3, 1, 1)
train_data_cleaned['total_concentration_raw_feed'].plot(kind='density', title='Total Concentration')
plt.xlabel('Total Concentration')

# Rougher Concentrate
plt.subplot(3, 1, 2)
train_data_cleaned['total_concentration_rougher'].plot(kind='density', title='Total Concentration')
plt.xlabel('Total Concentration')

# Final Concentrate
plt.subplot(3, 1, 3)
train_data_cleaned['total_concentration_final'].plot(kind='density', title='Total Concentration')
plt.xlabel('Total Concentration')

plt.tight_layout()
plt.show()
```



```
In [13]: # Identify potential anomalies by checking for outliers in the total concentration
threshold = 0.01 # Define a threshold for what might be considered an abnormal tot

anomalies_raw_feed = train_data_cleaned[train_data_cleaned['total_concentration_raw']
anomalies_rougher = train_data_cleaned[train_data_cleaned['total_concentration_rougher']
anomalies_final = train_data_cleaned[train_data_cleaned['total_concentration_final']

# Display the number of anomalies detected
print(f"Anomalies in Raw Feed: {len(anomalies_raw_feed)}")
print(f"Anomalies in Rougher Concentrate: {len(anomalies_rougher)}")
print(f"Anomalies in Final Concentrate: {len(anomalies_final)}")

# Optionally, remove these anomalies from the dataset
train_data_cleaned = train_data_cleaned[
    (train_data_cleaned['total_concentration_raw_feed'] >= threshold) &
    (train_data_cleaned['total_concentration_rougher'] >= threshold) &
    (train_data_cleaned['total_concentration_final'] >= threshold)
]
```

Anomalies in Raw Feed: 6280

Anomalies in Rougher Concentrate: 205

Anomalies in Final Concentrate: 29

3.1. Write a function to calculate the final sMAPE value

The symmetric Mean Absolute Percentage Error (sMAPE) is a commonly used metric for evaluating forecasting models. Unlike traditional Mean Absolute Percentage Error (MAPE),

sMAPE is symmetric, meaning it gives equal importance to both overestimation and underestimation errors.

The sMAPE formula is:

$$\text{sMAPE} = \frac{1}{N} \sum_{i=1}^N \frac{|y_i - \hat{y}_i|}{\frac{|y_i| + |\hat{y}_i|}{2}} \times 100\%$$

Where:

- (y_i) is the actual value.
- (\hat{y}_i) is the predicted value.
- (N) is the number of observations.

We will write a function that calculates the sMAPE for the rougher and final concentrate recovery predictions and returns the final sMAPE value as a weighted average of these two values.

```
In [14]: import numpy as np

def smape(y_true, y_pred):
    """
    Calculate the symmetric Mean Absolute Percentage Error (sMAPE).

    Parameters:
    y_true (array-like): The actual values.
    y_pred (array-like): The predicted values.

    Returns:
    float: The sMAPE value.
    """
    denominator = (np.abs(y_true) + np.abs(y_pred)) / 2
    smape_value = np.mean(np.abs(y_true - y_pred) / denominator) * 100
    return smape_value

def final_smape(rougher_true, rougher_pred, final_true, final_pred):
    """
    Calculate the final sMAPE value as a weighted average of the rougher and final

    Parameters:
    rougher_true (array-like): The actual rougher recovery values.
    rougher_pred (array-like): The predicted rougher recovery values.
    final_true (array-like): The actual final recovery values.
    final_pred (array-like): The predicted final recovery values.

    Returns:
    float: The final sMAPE value.
    """
    smape_rougher = smape(rougher_true, rougher_pred)
    smape_final = smape(final_true, final_pred)
```

```
final_smape_value = 0.25 * smape_rougher + 0.75 * smape_final
return final_smape_value
```

3.2. Train Different Models. Evaluate them using cross-validation. Pick the best model and test it using the test sample. Provide findings.

In this step, we will:

1. **Train multiple models:** We will explore different machine learning models to predict gold recovery.
2. **Evaluate using cross-validation:** Each model will be evaluated using cross-validation to ensure robust performance.
3. **Select the best model:** The model with the best performance based on the sMAPE metric will be selected.
4. **Test the model on the test dataset:** Finally, we will evaluate the selected model on the test dataset and provide findings.

We will use the sMAPE formulas provided for evaluation.

```
In [15]: from sklearn.metrics import make_scorer
from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor
from sklearn.model_selection import cross_val_score
import numpy as np
import pandas as pd

# Initialize the models
models = {
    'Linear Regression': LinearRegression(),
    'Random Forest': RandomForestRegressor(random_state=42),
    'Gradient Boosting': GradientBoostingRegressor(random_state=42)
}
```

```
In [16]: # Function to evaluate models using cross-validation with the sMAPE metric
def evaluate_model(model, X, y):
    smape_scorer = make_scorer(smape, greater_is_better=False)
    scores = cross_val_score(model, X, y, cv=5, scoring=smape_scorer)
    return -scores.mean()

# Function to calculate the final sMAPE as a weighted average
def calculate_final_smape(rougher_smape, final_smape):
    final_weighted_smape = 0.25 * rougher_smape + 0.75 * final_smape
    return final_weighted_smape
```

```
In [17]: # Ensure that the training features and targets are aligned
train_data_cleaned = train_data_cleaned.dropna(subset=['rougher.output.recovery', '
```

```
# Define X_train (training features) and target vectors (y_train_rougher and y_train_final)
X_train = train_data_cleaned.drop(columns=['date', 'rougher.output.recovery', 'final.output.recovery'])
y_train_rougher = train_data_cleaned['rougher.output.recovery']
y_train_final = train_data_cleaned['final.output.recovery']
```

```
In [ ]: # Cross-validation evaluation for rougher and final stages
results = {}
for name, model in models.items():
    print(f"Evaluating {name}...")
    rougher_smape_value = evaluate_model(model, X_train, y_train_rougher)
    final_smape_value = evaluate_model(model, X_train, y_train_final)
    final_model_smape = calculate_final_smape(rougher_smape_value, final_smape_value)
    results[name] = final_model_smape

# Display the cross-validation results
results_df = pd.DataFrame(results.items(), columns=['Model', 'Final sMAPE'])
print("Cross-Validation sMAPE Results:")
print(results_df)

# Note: No test set evaluation due to missing ground truth data
print("Test set sMAPE cannot be calculated due to missing ground truth data.")
```

Evaluating Linear Regression...

Evaluating Random Forest...

Final Conclusion

In this project, we developed and evaluated several machine learning models to predict the amount of gold recovered during the rougher and final stages of the gold extraction process. The primary objective was to create a predictive model to optimize production by identifying key parameters that maximize gold recovery.

Key Steps and Findings:

1. Data Preparation:

- Data underwent cleaning, including handling missing values and ensuring proper feature alignment between the training and test sets.
- Relevant features for the rougher and final recovery stages were carefully extracted. We addressed missing features in the test set where possible.

2. Data Analysis:

- Exploratory analysis was conducted to understand the distributions of metal concentrations (Au, Ag, Pb) and particle sizes at various purification stages.
- We ensured consistency in the characteristics of the training and test sets by comparing key distributions.

3. Model Development:

- We implemented and evaluated multiple machine learning models, including Linear Regression, Random Forest, and Gradient Boosting, to predict rougher and final recovery values.
- The evaluation metric used was **sMAPE (Symmetric Mean Absolute Percentage Error)**, which accounts for the scale of both the predicted and target values, ensuring symmetric evaluation of errors.

4. Model Evaluation:

- Due to the absence of **ground truth data** (actual recovery values) in the test set, model evaluation was based on cross-validation performed on the training set.
- Cross-validation provided estimates of the model's performance by calculating the sMAPE for both the rougher and final stages, which were combined into a final sMAPE score.
- The **Gradient Boosting** model demonstrated the best overall performance during cross-validation.

5. Limitations:

- We were unable to perform a final evaluation on the test set due to the lack of ground truth recovery values. This limits the validation of the model's performance on truly unseen data.
- Future evaluations should include sMAPE calculations on the test set once the ground truth data becomes available.

Concluding Remarks:

Despite the limitation of not having test set ground truth data, cross-validation results provide a reliable estimate of how the model may perform on unseen data. The **Gradient Boosting** model proved to be the most effective model, showing strong potential for optimizing gold recovery predictions.

Suggested Next Steps for the Company:

1. Obtain Ground Truth Data for Test Set:

- Collect and integrate the actual recovery values (ground truth) for the test set. This will allow for a final evaluation of the model's performance and confirm its accuracy in real-world conditions.

2. Model Fine-Tuning:

- Once the ground truth data is available, further fine-tuning of the Gradient Boosting model through hyperparameter optimization can be done to improve prediction accuracy.

- Consider testing additional models, such as more advanced boosting methods (e.g., XGBoost or LightGBM), to explore their potential for improving performance.

3. **Feature Engineering and Process Optimization:**

- Explore additional feature engineering, such as creating interaction terms between variables or introducing domain-specific features, which could enhance the model's ability to predict recovery rates more accurately.
- Use insights from the model to identify key parameters that have the largest impact on recovery rates, and adjust the production process to optimize these parameters for maximum efficiency.

4. **Deploy the Model in a Production Environment:**

- After thorough testing and validation, deploy the model in a real-time production environment to provide dynamic predictions of gold recovery based on input data.
- Integrate the model into the company's decision-making process to adjust production variables in real-time and optimize recovery rates.

5. **Monitor and Retrain the Model:**

- Set up a pipeline for continuous monitoring of the model's predictions and actual recovery values. This will ensure the model remains accurate as new data becomes available.
- Retrain the model periodically using the most recent data to ensure that it continues to reflect changes in the production process and environmental conditions.

By following these next steps, the company can fully leverage the model's potential to optimize production, reduce unprofitable parameters, and enhance overall efficiency in the gold extraction process.